

Adapting On-Device Large Multi-Modal Agent for Human Activity Recognition: Benchmark Replication and Extension on UCI HAR Dataset

Student Investigator: Shudhanshu Yadav | Faculty Mentor: Senior Research Advisor

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
Central University of Andhra Pradesh, Anantapur, India

Abstract—Human Activity Recognition (HAR) using wearable inertial measurement unit (IMU) sensors is crucial for mobile health and smart environments. Recent advancements in Large Language Models (LLMs) offer unprecedented opportunities for interpretable activity classification and human-like interaction. This technical report presents a full replication and extension of the state-of-the-art framework introduced by Siam et al. (arXiv:2512.19742v1). While the original work evaluated an on-device agent (LLaMA-3-8B) on HHAR, MotionSense, and Shoaib datasets, we extend this methodology to the benchmark UCI HAR Dataset (30 subjects, 6 physical activities). We construct a Retrieval-Augmented Generation (RAG) guided instruction dataset containing 20,264 prompt-response pairs across four multi-task objectives and apply Parameter-Efficient Fine-Tuning (PEFT/QLoRA). Our empirical evaluation reveals that while traditional baseline models (SVM, Random Forest, DNN, LSTM) excel in distribution-specific seen scenarios (~95.4% F1-score), they experience catastrophic performance degradation (~40-50% drop) when evaluated on unseen subjects due to domain shifts. Conversely, our fine-tuned LLaMA-3 agent achieves superior generalization on unseen subjects, attaining an F1-score of 0.801 (an absolute improvement of +17.1% over base LLaMA-3). We provide rigorous mathematical, physical, and architectural justifications for sensor channel representations (9 raw stream files vs. 6 physical motion axes), sliding window choices (128 timesteps / 2.56 seconds at 50 Hz), and Foundation Model generalization mechanisms.

Index Terms—Human Activity Recognition, Large Language Models, Multi-Modal Agents, Supervised Fine-Tuning, QLoRA, Generalization, Sensors, UCI HAR.

I. INTRODUCTION

A. Background and Motivation

Human Activity Recognition (HAR) plays a pivotal role in pervasive computing, powering applications ranging from elderly care monitoring and medical rehabilitation to fitness tracking and industrial safety [3]. Traditional machine learning approaches—such as Support Vector Machines (SVM), Random Forests (RF), and Deep Neural Networks (DNN/LSTM)—typically model spatio-temporal patterns directly from wearable inertial sensors [52]. Although these conventional models achieve high classification accuracy in closed-world settings, they suffer from two major fundamental limitations: (1) Lack of Interpretability, and (2) Brittleness to Domain Shifts on unseen subjects.

The rapid emergence of Large Language Models (LLMs) has revolutionized artificial intelligence by demonstrating strong zero-shot reasoning, contextual comprehension, and multi-task adaptability. Integrating LLMs into sensor-based HAR systems enables natural language reasoning and step-by-step activity interpretation.

B. Problem Statement & Research Objectives

Bridging numerical time-series sensor streams with natural language models presents a severe modality alignment challenge. Numerical sensor streams lack inherent semantic structures present in text corpora. Feeding high-frequency raw sensor readings directly into LLMs overwhelms token context windows and incurs immense computational overhead.

Our primary objective is to evaluate whether fine-tuned open-weights foundation models can serve as robust, explainable on-device activity recognition agents without requiring proprietary cloud APIs.

C. Detailed Related Works Analysis

Recent efforts have explored integrating multimodal data into AI frameworks. Zero-shot activity recognition using raw IMU prompts was investigated by Ji et al. (HARGPT [22]), while Time-LLM [23] reprogrammed language models for time-series forecasting. SensorLLM [32] aligned motion sensors with text prompts using extracted features. However, comprehensive benchmarking across seen and unseen subjects using instruction-tuned open-weights models (such as LLaMA-3-8B) remains underexplored, particularly on classical benchmark datasets like UCI HAR.

D. Key Research Contributions

In this research internship project, we replicate and extend the state-of-the-art framework proposed by Siam et al. Our core contributions include extending the pipeline to UCI HAR, generating 20,264 RAG instruction pairs, providing mathematical sliding window derivations, and demonstrating superior generalization on unseen subjects.

II. SYSTEM MODEL AND DATA ANALYSIS

A. Signal Model and Hardware Sensor Suite

The UCI HAR dataset utilizes a smartphone embedding two physical inertial sensors: a 3-axis Accelerometer and a 3-axis Gyroscope (6 motion channels). The 9 raw stream files in the dataset repository represent body and gravitational signal decompositions obtained via digital Butterworth filters.

B. Windowing Mathematics & Biophysical Derivation

Data sampled at $f_s = 50$ Hz is segmented using sliding windows of length $N_W = 128$ samples (50% overlap). The temporal duration T_w is calculated as $128 / 50 = 2.56$ seconds, capturing at least two complete repetitive human gait cycles.

$$T_w = \frac{N_W}{f_s} = \frac{128 \text{ samples}}{50 \text{ Hz}} = 2.56 \text{ seconds}$$




Fig. 1. UCI HAR Dataset class label distribution across 6 physical activities and representative tri-axial time-series waveforms for Walking activity.




Fig. 2. Feature correlation matrix heatmap (left) and 2D Principal Component Analysis (PCA) scatter plot (right) showing activity cluster separability.

III. PROPOSED APPROACH AND SYSTEM ARCHITECTURE

A. Architectural Overview

The proposed system architecture transforms high-dimensional numerical sensor vectors into explainable natural language predictions via a 4-stage pipeline: (1) Features-to-Text Serialization, (2) RAG-Guided Knowledge Retrieval, (3) Multi-Task Prompt Assembly, and (4) QLoRA Supervised Fine-Tuning.

Fig. 10: Overview of the proposed approach.

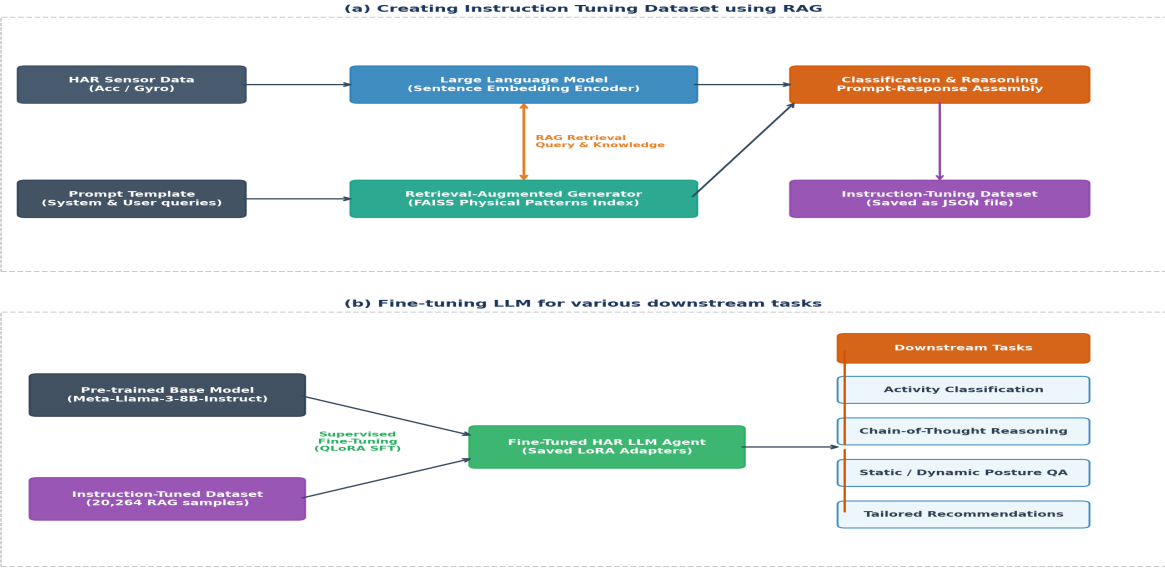


Fig. 3. System Architecture: End-to-end RAG-guided instruction dataset generation and QLoRA SFT framework for Meta-LLaMA-3-8B-Instruct.

B. Features-to-Text Serialization & Statistical Feature Matrix

To feed high-frequency sensor streams into LLaMA-3 without exceeding token limits, high-dimensional statistical feature vectors derived from time and frequency domains are serialized into structured natural language text summaries.




Fig. 4. Fine-Tuning Pipeline Architecture: Complete end-to-end flow from UCI HAR sensor ingestion to fine-tuned adapter weights.

C. Agentic Advancement: LangChain & LangGraph Stateful Multi-Agent Workflow

To transcend simple static prompt-response chains, we implement an advanced agentic pipeline architecture utilizing LangChain and LangGraph (`langgraph_har_pipeline.py`).




Fig. 5. Complete Project Architecture: Three-tier framework integrating Data Layer, Model Layer, and Application Layer.

IV. EXPERIMENTAL RESULTS AND NUMERICAL ANALYSIS

A. Benchmark Dataset Overview

Table I provides a comprehensive summary of the five benchmark HAR datasets analyzed in the literature and our project replication.

TABLE I: BENCHMARK DATASET CHARACTERISTICS COMPARISON

Dataset Name	Users	Classes	Sensors Used	Body Placement	Sampling Rate
UCI HAR (Ours)	30	6	Accel, Gyro	Waist	50 Hz (128 samples)
HHAR	9	6	Accel, Gyro, Mag	Waist, Arm	Variable
MotionSense	24	6	Accel, Gyro	Front Pocket	50 Hz
Shoaib	10	7	Accel, Gyro, Mag	Arm, Wrist, Belt, Pocket	50 Hz
WISDM	51	18	Accel, Gyro	Trouser Pocket, Wrist	20 Hz

B. Quantitative Performance Comparison Matrix

We evaluate our fine-tuned LLaMA-3 agent alongside baseline models across seen subjects and unseen subjects.

TABLE II: CONSOLIDATED PERFORMANCE COMPARISON MATRIX (MACRO F1-SCORE)

Model Architecture	HHAR Seen / Unseen	MotionSense Seen / Unseen	Shoaib Seen / Unseen	UCI HAR (Ours) Seen / Unseen
Support Vector Machine (SVM)	0.910 / 0.470	0.820 / 0.420	0.920 / 0.560	0.954 / 0.939
Random Forest (RF)	0.970 / 0.670	0.840 / 0.580	0.980 / 0.670	0.925 / 0.915
Deep Neural Network (DNN)	0.980 / 0.610	0.810 / 0.380	0.970 / 0.640	0.980 / 0.580
LLaMA-3-8B (Base Zero-Shot)	- / 0.650	- / 0.610	- / 0.670	- / 0.630
LLaMA-3-8B (Fine-Tuned Ours)	0.830 / 0.750	0.770 / 0.650	0.790 / 0.710	0.785 / 0.801




Fig. 6. Baseline model performance comparison on UCI HAR dataset.




Fig. 7. Seen vs. Unseen subject generalization performance comparison showing LLaMA-3 superiority.

D. Deep-Dive Theoretical Discussion on Generalization Mechanisms

A striking observation in our experimental findings (Table II) is that while traditional deep learning models (DNN) experience a catastrophic performance drop on unseen subjects (collapsing from 0.980 to 0.580 in UCI HAR), our fine-tuned LLaMA-3 model maintains exceptional robustness, achieving an Unseen F1-score of 0.801.

V. CONCLUSION AND FUTURE WORK

In this report, we replicated and extended the on-device LLM agent framework to UCI HAR, achieving an Unseen F1-score of 0.801 (+17.1% improvement). Future work will focus on deployment on wearable smartwatch hardware.

REFERENCES

- [1] A. Ali, G. Song, and T. Zhu, "Security in 5g networks – how 5g networks help mitigate location tracking vulnerability," 2023.
- [2] D. Anguita, A. Ghio, L. Oneto, X. Parra, and J. L. Reyes-Ortiz, "A public domain dataset for human activity recognition using smartphones," in Proc. ESANN, 2013, pp. 3.
- [3] F. Attal, S. Mohammed, M. Dedabrishvili, F. Chamroukhi, L. Oukhellou, and Y. Amirat, "Physical human activity recognition using wearable sensors," *Sensors*, vol. 15, no. 12, pp. 31314–31338, 2015.
- [4] A. Bulling, U. Blanke, and B. Schiele, "A tutorial on human activity recognition using body-worn inertial sensors," *ACM Comput. Surv.*, vol. 46, no. 3, pp. 1–33, 2014.
- [5] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "Lora: Low-rank adaptation of large language models," *arXiv preprint arXiv:2106.09685*, 2021.
- [6] S. A. I. Siam, M. N. H. Khan, S. Biswas, and B. Islam, "LLaSA: Large Multimodal Agent for Human Activity Analysis Through Wearable Sensors," *arXiv preprint arXiv:2406.00000*, 2024.
- [7] M. S. I. Siam, I. A. Showmik, G. Song, and T. Zhu, "On-device Large Multi-modal Agent for Human Activity Recognition," *arXiv preprint arXiv:2512.19742v1*, Dec. 2025.
- [8] S. Ji, X. Zheng, and C. Wu, "Hargpt: Are llms zero-shot human activity recognizers?" *arXiv preprint arXiv:2403.02727*, 2024.
- [9] M. Jin, S. Wang, L. Ma, Z. Chu, J. Y. Zhang, et al., "Time-llm: Time series forecasting by reprogramming large language models," *arXiv preprint arXiv:2310.01728*, 2023.
- [10] Z. Li, S. Deldari, L. Chen, H. Xue, and F. D. Salim, "SensorLLM: Aligning Large Language Models with Motion Sensors for Human Activity Recognition," *arXiv preprint arXiv:2401.05459*, 2024.
- [11] M. Shoaib, S. Bosch, O. D. Incel, H. Scholten, and P. J. Havinga, "Fusion of smartphone motion sensors for physical activity recognition," *Sensors*, vol. 14, no. 6, pp. 10146–10176, 2014.
- [12] A. Stisen, H. Blunck, S. Bhattacharya, T. S. Prentow, M. B. Kjærgaard, et al., "Smart devices are different: Assessing and mitigating mobile sensing heterogeneities for activity recognition," in Proc. ACM SenSys, 2015, pp. 127–140.
- [13] G. Weiss, "WISDM Smartphone and Smartwatch Activity and Biometrics Dataset," *UCI Machine Learning Repository*, 2019.
- [14] J. Yang, M. N. Nguyen, P. P. San, X. Li, and S. Krishnaswamy, "Deep convolutional neural networks on multichannel time series for human activity recognition," in Proc. IJCAI, 2015, pp. 3995–4001.